

# Evaluation of Pointer Swizzling Techniques for DBMS Buffer Management

Max Gilbert

University of Kaiserslautern

*m\_gilbert13@cs.uni-kl.de*

June 8, 2017

# Table of Contents

## Pointer Swizzling as in “In-Memory Performance for Big Data”

- Locate Pages in the Buffer Pool without Pointer Swizzling

- Locate Pages in the Buffer Pool with Pointer Swizzling

## Performance Evaluation of the Buffer Management Utilizing Pointer Swizzling

- Expected Performance

- Measured Performance

- Conclusion

## Page Eviction Strategies in the Context of Pointer Swizzling

- Probable pitfalls when Implementing a Page Eviction Strategy for a DBMS Buffer Manager

- Evaluated Page Replacement Strategies

- Performance Evaluation

- Conclusion

## Section 1

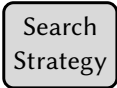
# Pointer Swizzling as in “In-Memory Performance for Big Data”

## Subsection 1

### Locate Pages in the Buffer Pool without Pointer Swizzling


# Overview of Search Strategies ([HR01])

# Overview of Search Strategies ([HR01])




Search  
Strategy

## Overview of Search Strategies ([HR01])

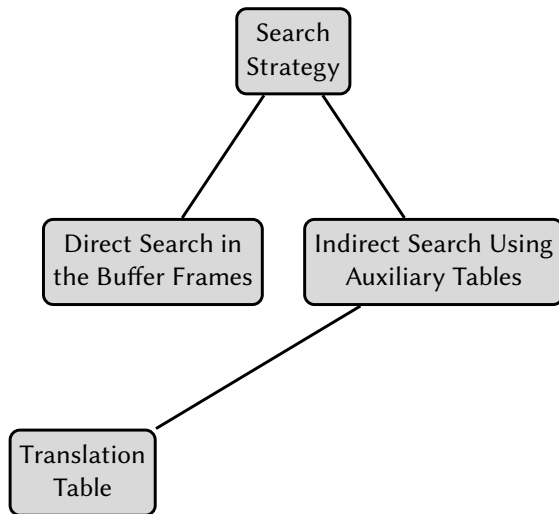


## Overview of Search Strategies ([HR01])






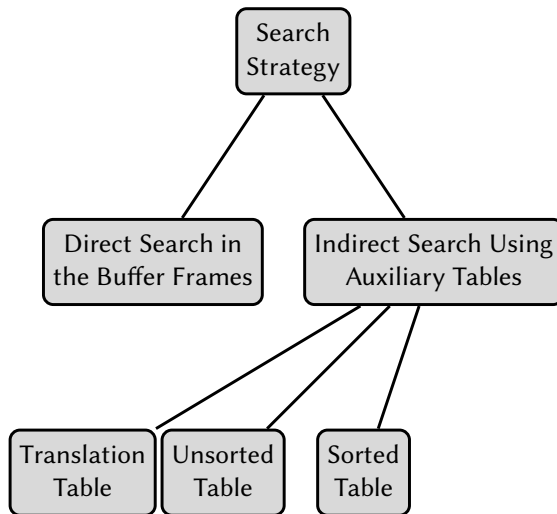
## Overview of Search Strategies ([HR01])




## Overview of Search Strategies ([HR01])



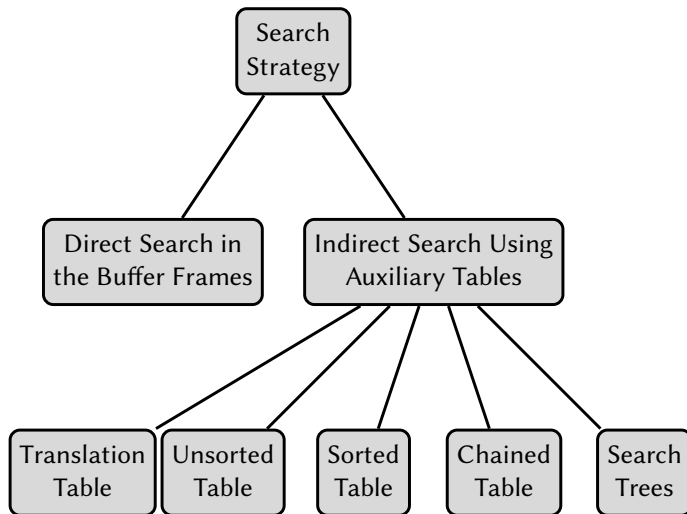
## Overview of Search Strategies ([HR01])



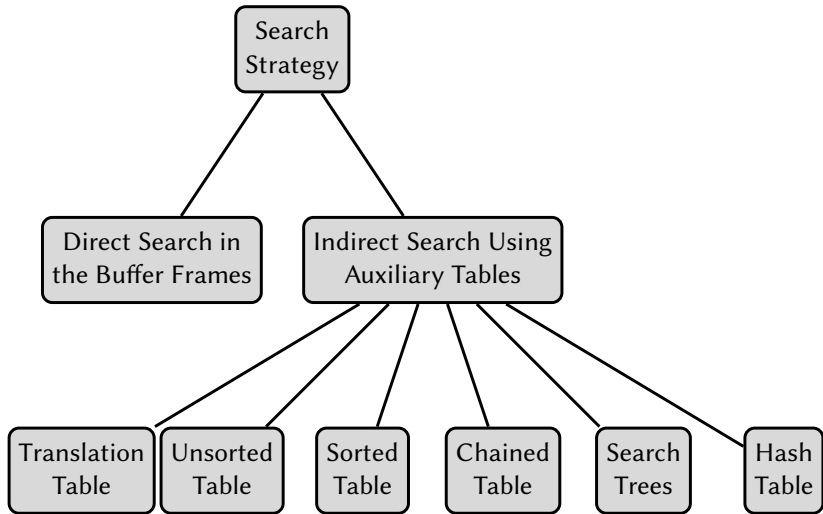
# Overview of Search Strategies ([HR01])



## Overview of Search Strategies ([HR01])



## Overview of Search Strategies ([HR01])



# Direct Search in the Buffer Frames & Unsorted Table

# Direct Search in the Buffer Frames & Unsorted Table

## Direct Search in the Buffer Frames



# Direct Search in the Buffer Frames & Unsorted Table

## Direct Search in the Buffer Frames

- Checks in each buffer frame the page ID of the contained page

# Direct Search in the Buffer Frames & Unsorted Table

## Direct Search in the Buffer Frames

- ▶ Checks in each buffer frame the page ID of the contained page
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}\left(\frac{n}{2}\right)$ ,  $T_{\text{worst}}^{\text{search}} \in \mathcal{O}(n)$

# Direct Search in the Buffer Frames & Unsorted Table

## Direct Search in the Buffer Frames

- ▶ Checks in each buffer frame the page ID of the contained page
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}\left(\frac{n}{2}\right)$ ,  $T_{\text{worst}}^{\text{search}} \in \mathcal{O}(n)$
- ▶ The usage of virtual memory management can result in extensive swapping due to read access to many pages!

# Direct Search in the Buffer Frames & Unsorted Table

## Direct Search in the Buffer Frames

- ▶ Checks in each buffer frame the page ID of the contained page
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}\left(\frac{n}{2}\right)$ ,  $T_{\text{worst}}^{\text{search}} \in \mathcal{O}(n)$
- ▶ The usage of virtual memory management can result in extensive swapping due to read access to many pages!

## Unsorted Table

# Direct Search in the Buffer Frames & Unsorted Table

## Direct Search in the Buffer Frames

- ▶ Checks in each buffer frame the page ID of the contained page
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}\left(\frac{n}{2}\right)$ ,  $T_{\text{worst}}^{\text{search}} \in \mathcal{O}(n)$
- ▶ The usage of virtual memory management can result in extensive swapping due to read access to many pages!

## Unsorted Table

- ▶ Auxiliary data structure of size  $S_{\text{pace}} \in \mathcal{O}(n)$

0	1	2	3	4	5	6	7	8
7785	6977	4347	3380	5610	6376	4877	3332	3354

**Figure:** An unsorted table used to map buffer frames to page IDs.

# Direct Search in the Buffer Frames & Unsorted Table

## Direct Search in the Buffer Frames

- ▶ Checks in each buffer frame the page ID of the contained page
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}\left(\frac{n}{2}\right)$ ,  $T_{\text{worst}}^{\text{search}} \in \mathcal{O}(n)$
- ▶ The usage of virtual memory management can result in extensive swapping due to read access to many pages!

## Unsorted Table

- ▶ Auxiliary data structure of size  $S_{\text{page}} \in \mathcal{O}(n)$
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}\left(\frac{n}{2}\right)$ ,  $T_{\text{worst}}^{\text{search}} \in \mathcal{O}(n)$

0	1	2	3	4	5	6	7	8
7785	6977	4347	3380	5610	6376	4877	3332	3354

**Figure:** An unsorted table used to map buffer frames to page IDs.

# Translation Table

## Translation Table

- ▶ Auxiliary data structure with one entry per page in the database  
 $\implies S_{\text{pace}} \in \mathcal{O}(p)$

[illegible]

**Figure:** A translation table used to map page IDs to buffer frames.



# Translation Table

- ▶ Auxiliary data structure with one entry per page in the database  
 $\implies S_{\text{pace}} \in \mathcal{O}(p)$
- ▶  $T^{\text{search}} \in \mathcal{O}(1), T^{\text{insert}} \in \mathcal{O}(1)$

0	.	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
⋮	⋮	3352	.	3378	.	4345	.	4875	.	5608	.	6374	.	6975	.	7783	.
3331	.	3353	.	3379	.	4346	.	4876	.	5609	.	6375	.	6976	.	7784	.
3332	7	3354	8	3380	3	4347	2	4877	6	5610	4	6376	5	6977	1	7785	0
3333	.	3355	.	3381	.	4348	.	4878	.	5611	.	6377	.	6978	.	7786	.
⋮	⋮	3356	.	3382	.	4349	.	4879	.	5612	.	6378	.	6979	.	7787	.
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

**Figure:** A translation table used to map page IDs to buffer frames.

# Sorted & Chained Table

# Sorted & Chained Table

## Sorted Table

# Sorted & Chained Table

## Sorted Table

- Auxiliary data structure using a table sorted by page ID only containing cached pages

3332	3354	3380	4347	4877	5610	6376	6977	7785
→ 7	→ 8	→ 3	→ 2	→ 6	→ 4	→ 5	→ 1	→ 0

**Figure:** A sorted table used to map page IDs to buffer frames.

# Sorted & Chained Table

## Sorted Table

- ▶ Auxiliary data structure using a table sorted by page ID only containing cached pages
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log_2 n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(n \log_2 n)$

3332	3354	3380	4347	4877	5610	6376	6977	7785
→ 7	→ 8	→ 3	→ 2	→ 6	→ 4	→ 5	→ 1	→ 0

**Figure:** A sorted table used to map page IDs to buffer frames.

# Sorted & Chained Table

## Sorted Table

- ▶ Auxiliary data structure using a table sorted by page ID only containing cached pages
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log_2 n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(n \log_2 n)$

3332	3354	3380	4347	4877	5610	6376	6977	7785
→ 7	→ 8	→ 3	→ 2	→ 6	→ 4	→ 5	→ 1	→ 0

**Figure:** A sorted table used to map page IDs to buffer frames.

## Chained Table

# Sorted & Chained Table

## Sorted Table

- ▶ Auxiliary data structure using a table sorted by page ID only containing cached pages
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log_2 n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(n \log_2 n)$

3332	3354	3380	4347	4877	5610	6376	6977	7785
→ 7	→ 8	→ 3	→ 2	→ 6	→ 4	→ 5	→ 1	→ 0

Figure: A sorted table used to map page IDs to buffer frames.

## Chained Table

- ▶ Auxiliary data structure using a linked list sorted by page ID only containing cached pages




Figure: A chained table used to map page IDs to buffer frames.

# Sorted & Chained Table

## Sorted Table

- ▶ Auxiliary data structure using a table sorted by page ID only containing cached pages
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log_2 n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(n \log_2 n)$

3332	3354	3380	4347	4877	5610	6376	6977	7785
→ 7	→ 8	→ 3	→ 2	→ 6	→ 4	→ 5	→ 1	→ 0

Figure: A sorted table used to map page IDs to buffer frames.

## Chained Table

- ▶ Auxiliary data structure using a linked list sorted by page ID only containing cached pages
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log_2 n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(\log_2 n)$




Figure: A chained table used to map page IDs to buffer frames.



# Sorted & Chained Table

## Sorted Table

- ▶ Auxiliary data structure using a table sorted by page ID only containing cached pages
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log_2 n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(n \log_2 n)$

3332	3354	3380	4347	4877	5610	6376	6977	7785
→ 7	→ 8	→ 3	→ 2	→ 6	→ 4	→ 5	→ 1	→ 0

Figure: A sorted table used to map page IDs to buffer frames.

## Chained Table

- ▶ Auxiliary data structure using a linked list sorted by page ID only containing cached pages
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log_2 n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(\log_2 n)$
- ▶ Binary search requires more links!





Figure: A chained table used to map page IDs to buffer frames.

# Search Trees

## Search Trees


- Auxiliary data structure is similar to the one of the chained table



**Figure:** A balanced search tree used to map page IDs to buffer frames.

## Search Trees


- ▶ Auxiliary data structure is similar to the one of the chained table
- ▶ Many different data structures like AVL-trees, red-black trees or splay trees can be used



**Figure:** A balanced search tree used to map page IDs to buffer frames.

## Search Trees


- ▶ Auxiliary data structure is similar to the one of the chained table
- ▶ Many different data structures like AVL-trees, red-black trees or splay trees can be used
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(\log n)$



**Figure:** A balanced search tree used to map page IDs to buffer frames.

## Search Trees


- ▶ Auxiliary data structure is similar to the one of the chained table
- ▶ Many different data structures like AVL-trees, red-black trees or splay trees can be used
- ▶  $T_{\text{avg}}^{\text{search}} \in \mathcal{O}(\log n)$ ,  $T_{\text{avg}}^{\text{insert}} \in \mathcal{O}(\log n)$
- ▶ The worst case costs and the worst cases vary between the different search tree data structures



**Figure:** A balanced search tree used to map page IDs to buffer frames.


# Hash Table

# Hash Table






# Hash Table




- Each page ID is mapped to a hash bucket using a hash function

# Hash Table




- ▶ Each page ID is mapped to a hash bucket using a hash function
- ▶ Only the page IDs of buffered pages are in the hash table

# Hash Table



- ▶ Each page ID is mapped to a hash bucket using a hash function
- ▶ Only the page IDs of buffered pages are in the hash table
- ▶ If a hash bucket is full, a chained bucket gets added

# Hash Table



- ▶ Each page ID is mapped to a hash bucket using a hash function
- ▶ Only the page IDs of buffered pages are in the hash table
- ▶ If a hash bucket is full, a chained bucket gets added
- ▶  $T_{avg}^{search} \in \mathcal{O}(1)$ ,  
 $T_{avg}^{insert} \in \mathcal{O}(1)$ ,  
 $T_{worst}^{search} \in \mathcal{O}(n)$

## Locate Pages in Buffer Pool with Hash Table ([Gra+14])

# Locate Pages in Buffer Pool with Hash Table ([Gra+14])



Buffer pool  
page image


## Locate Pages in Buffer Pool with Hash Table ([Gra+14])



Buffer pool  
page image


Search key

## Locate Pages in Buffer Pool with Hash Table ([Gra+14])

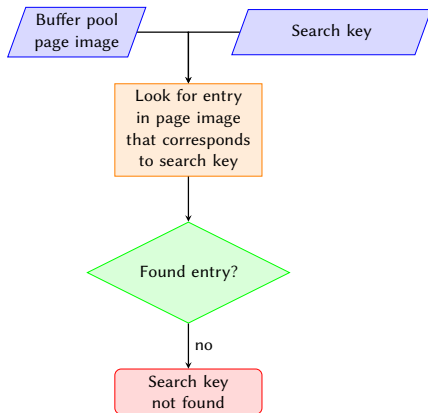




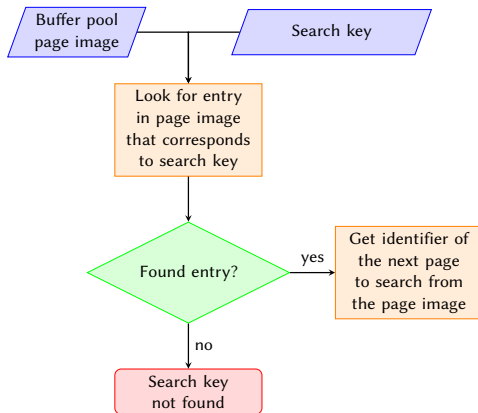
## Locate Pages in Buffer Pool with Hash Table ([Gra+14])




## Locate Pages in Buffer Pool with Hash Table ([Gra+14])



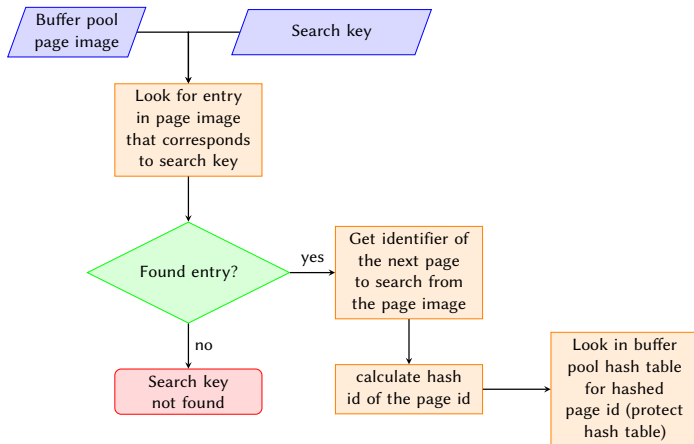
## Locate Pages in Buffer Pool with Hash Table ([Gra+14])



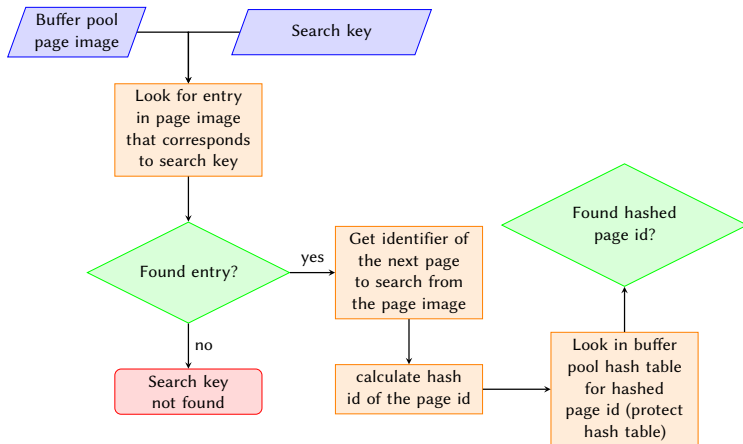
## Locate Pages in Buffer Pool with Hash Table ([Gra+14])



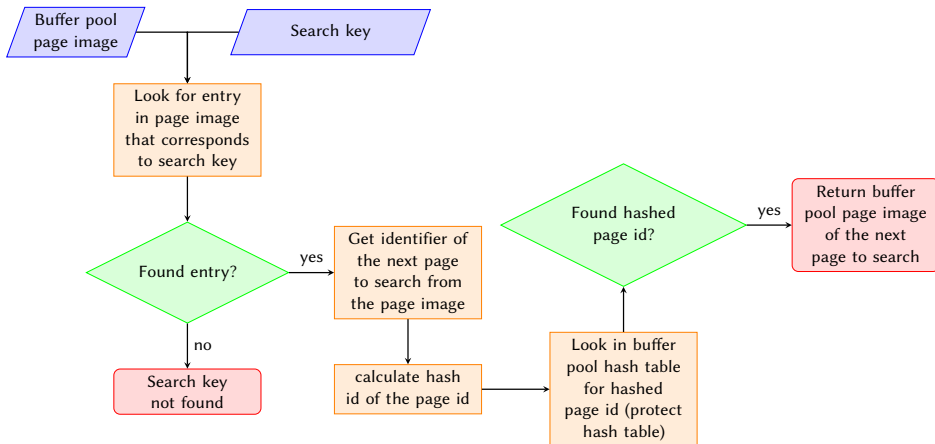
## Locate Pages in Buffer Pool with Hash Table ([Gra+14])



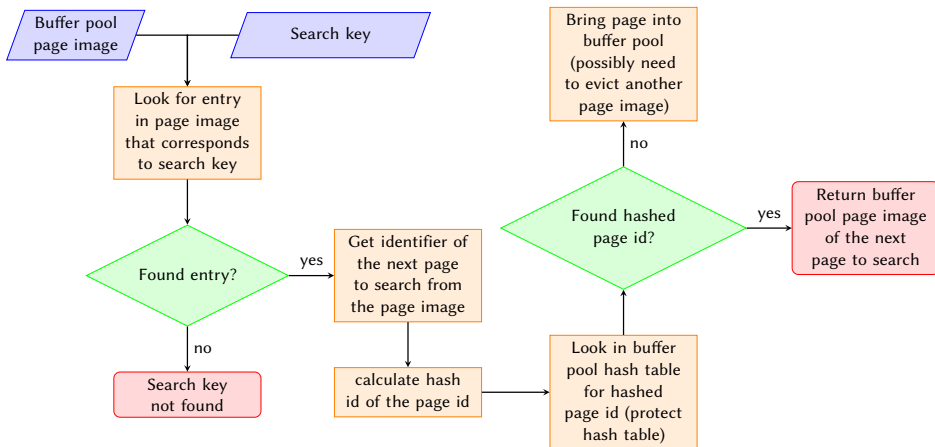
## Locate Pages in Buffer Pool with Hash Table ([Gra+14])



## Locate Pages in Buffer Pool with Hash Table ([Gra+14])

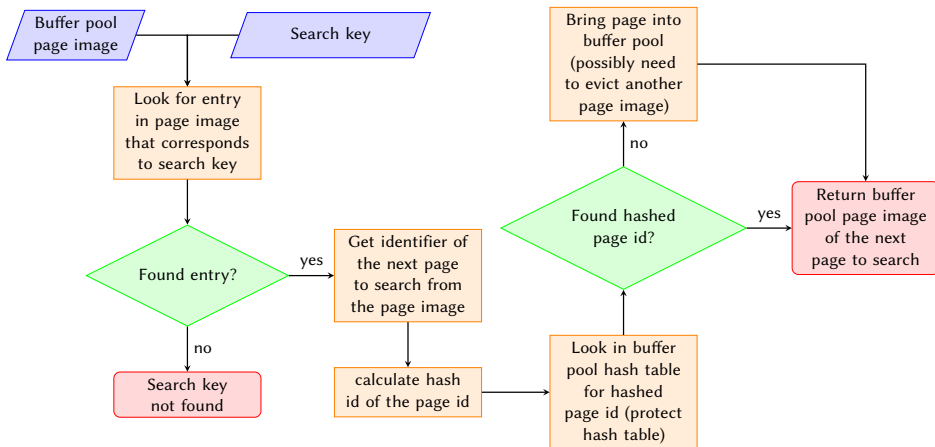


## Locate Pages in Buffer Pool with Hash Table ([Gra+14])





## Locate Pages in Buffer Pool with Hash Table ([Gra+14])



## Subsection 2

### Locate Pages in the Buffer Pool with Pointer Swizzling

# Pointer Swizzling


## Definition

To swizzle a pointer means to transform the address of the persistent object referenced there to a more direct address of the transient object in a way that this transformation could be used during multiple indirections of this pointer ([Mos92]).

# Classification of the Pointer Swizzling Approach following [WD95]


# Classification of the Pointer Swizzling Approach following [WD95]

no-swizzling




swizzling


# Classification of the Pointer Swizzling Approach following [WD95]



# Classification of the Pointer Swizzling Approach following [WD95]




# Classification of the Pointer Swizzling Approach following [WD95]







# Classification of the Pointer Swizzling Approach following [WD95]



# Classification of the Pointer Swizzling Approach following [WD95]



# Classification of the Pointer Swizzling Approach following [WD95]



# Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])

# Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])



Buffer pool  
page image


## Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])




Buffer pool  
page image

Search key

## Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])




## Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])







## Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])



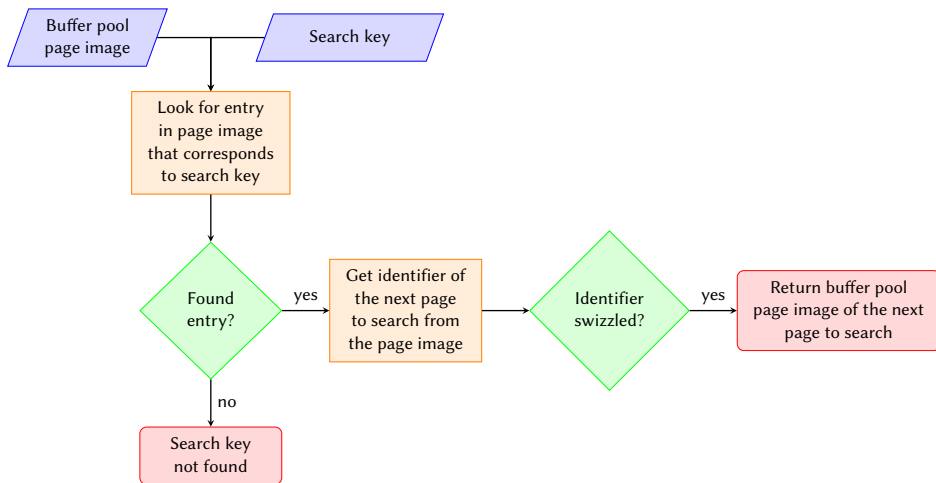
# Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])



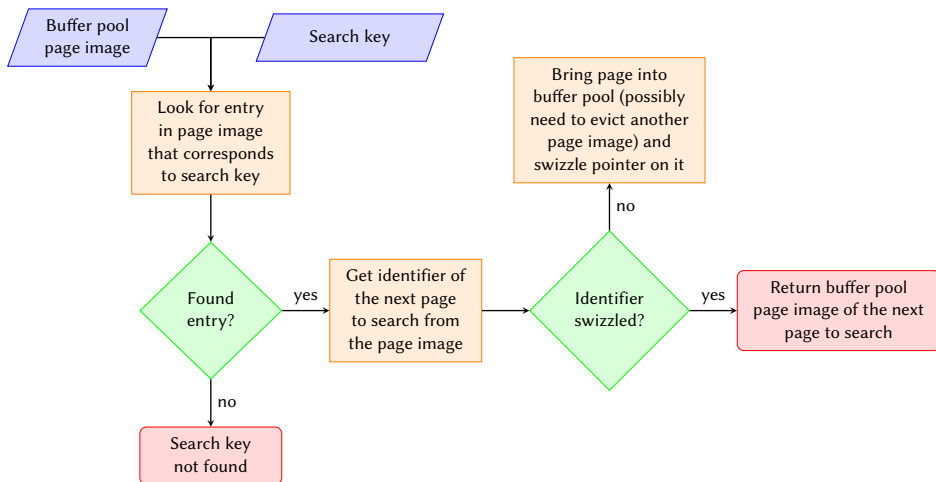
# Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])



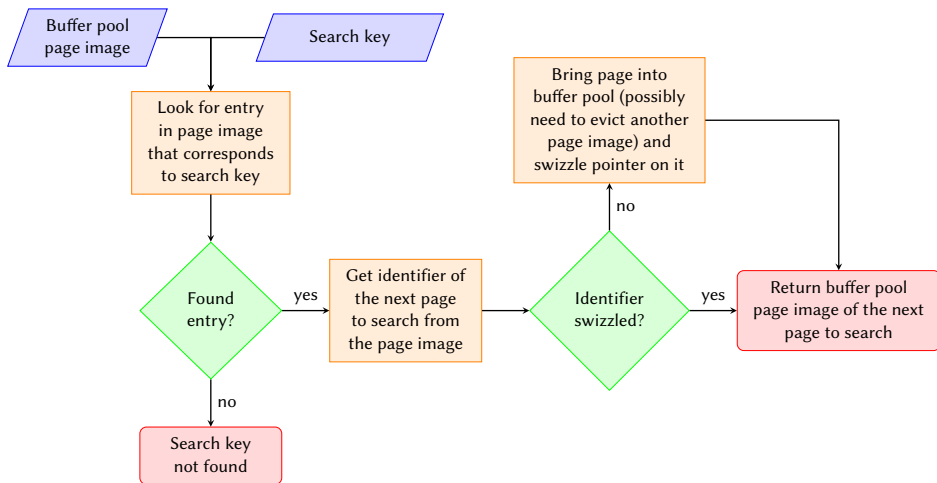
# Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])



# Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])



# Locate Pages in Buffer Pool w/ Pointer Swizzling ([Gra+14])



## Section 2


# Performance Evaluation of the Buffer Management Utilizing Pointer Swizzling

## Subsection 1


### Expected Performance




# Performance of Different Buffer Pool Sizes




# Performance of Different Buffer Pool Sizes




# Performance of Different Buffer Pool Sizes



## Performance of Different Buffer Pool Sizes




# Buffer Management with and without Pointer Swizzling




# Buffer Management with and without Pointer Swizzling



# Buffer Management with and without Pointer Swizzling




# Buffer Management with and without Pointer Swizzling

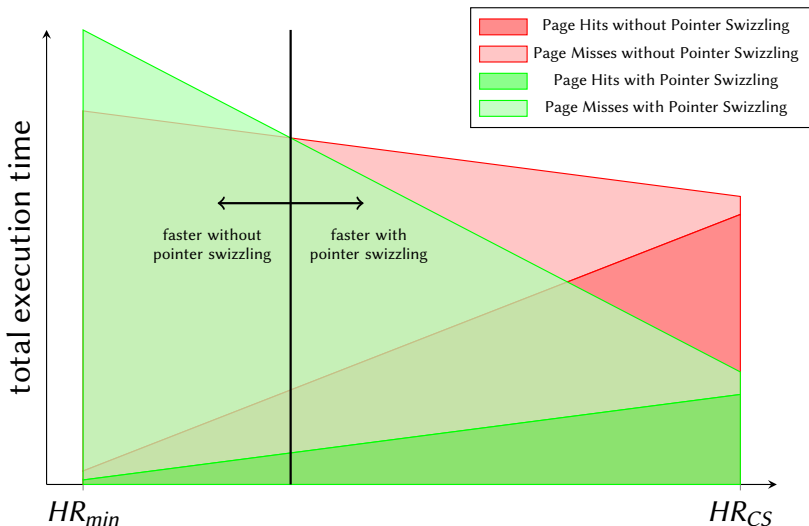




# Buffer Management with and without Pointer Swizzling




# Buffer Management with and without Pointer Swizzling




## Subsection 2

### Measured Performance


# Transaction Throughput (TPC-C)




# Transaction Throughput (TPC-C)




# Transaction Throughput (TPC-C)



# Transaction Throughput (TPC-B)




# Transaction Throughput (TPC-B)







# Transaction Throughput (TPC-B)




# Buffer Pool Performance Acquiring Shared Latches




# Buffer Pool Performance Acquiring Shared Latches




# Buffer Pool Performance Acquiring Shared Latches



# Buffer Pool Performance Acquiring Shared Latches



# Buffer Pool Performance Acquiring Shared Latches



## Subsection 3

## Conclusion

# Conclusion



# Conclusion

## Overall Performance

# Conclusion

## Overall Performance

- ▶ Pointer swizzling couldn't improve the performance on TPC-C benchmark runs with a duration of 10 min.

# Conclusion

## Overall Performance

- ▶ Pointer swizzling couldn't improve the performance on TPC-C benchmark runs with a duration of 10 min.
- ▶ The page hits after the cold start couldn't compensate the overhead of pointer swizzling during the cold start.

# Conclusion

## Overall Performance

- ▶ Pointer swizzling couldn't improve the performance on TPC-C benchmark runs with a duration of 10 min.
- ▶ The page hits after the cold start couldn't compensate the overhead of pointer swizzling during the cold start.
- ▶ A continuously running DB with large buffer pool could profit from pointer swizzling.

# Conclusion

## Overall Performance

- ▶ Pointer swizzling couldn't improve the performance on TPC-C benchmark runs with a duration of 10 min.
- ▶ The page hits after the cold start couldn't compensate the overhead of pointer swizzling during the cold start.
- ▶ A continuously running DB with large buffer pool could profit from pointer swizzling.

## Buffer Pool Performance

# Conclusion

## Overall Performance

- ▶ Pointer swizzling couldn't improve the performance on TPC-C benchmark runs with a duration of 10 min.
- ▶ The page hits after the cold start couldn't compensate the overhead of pointer swizzling during the cold start.
- ▶ A continuously running DB with large buffer pool could profit from pointer swizzling.

## Buffer Pool Performance

- ▶ A page hit is faster when pointer swizzling is activated.

# Conclusion

## Overall Performance

- ▶ Pointer swizzling couldn't improve the performance on TPC-C benchmark runs with a duration of 10 min.
- ▶ The page hits after the cold start couldn't compensate the overhead of pointer swizzling during the cold start.
- ▶ A continuously running DB with large buffer pool could profit from pointer swizzling.

## Buffer Pool Performance

- ▶ A page hit is faster when pointer swizzling is activated.
- ▶ A page miss is slower when pointer swizzling is activated.

# Conclusion

## Overall Performance

- ▶ Pointer swizzling couldn't improve the performance on TPC-C benchmark runs with a duration of 10 min.
- ▶ The page hits after the cold start couldn't compensate the overhead of pointer swizzling during the cold start.
- ▶ A continuously running DB with large buffer pool could profit from pointer swizzling.

## Buffer Pool Performance

- ▶ A page hit is faster when pointer swizzling is activated.
- ▶ A page miss is slower when pointer swizzling is activated.
- ▶ After the cold start phase, activated pointer swizzling will improve the buffer pool performance for large buffer pools.



## Section 3


# Page Eviction Strategies in the Context of Pointer Swizzling

# Motivation not to Analyze Different Page Eviction Strategies

# Motivation not to Analyze Different Page Eviction Strategies

- ▶ Even LRU results in decent hit rates

TPC-C with Warehouses: 100, Threads: 25



# But ...

## But ...

- ▶ Page reference pattern containing a loop slightly to long to fit in the buffer pool:

## But ...

- ▶ Page reference pattern containing a loop slightly too long to fit in the buffer pool:
  - ▶ **OPT**: Hit rate close to 1

## But ...

- ▶ Page reference pattern containing a loop slightly to long to fit in the buffer pool:
  - ▶ **OPT**: Hit rate close to 1
  - ▶ **LRU**: Hit rate of 0

## But ...

- ▶ Page reference pattern containing a loop slightly too long to fit in the buffer pool:
  - ▶ **OPT**: Hit rate close to 1
  - ▶ **LRU**: Hit rate of 0
- ▶ Some pages get referenced very frequently for a limited time:



## But ...

- ▶ Page reference pattern containing a loop slightly too long to fit in the buffer pool:
  - ▶ **OPT**: Hit rate close to 1
  - ▶ **LRU**: Hit rate of 0
- ▶ Some pages get referenced very frequently for a limited time:
  - ▶ **OPT**: Pages would be evicted after their last reference

## But ...

- ▶ Page reference pattern containing a loop slightly too long to fit in the buffer pool:
  - ▶ **OPT**: Hit rate close to 1
  - ▶ **LRU**: Hit rate of 0
- ▶ Some pages get referenced very frequently for a limited time:
  - ▶ **OPT**: Pages would be evicted after their last reference
  - ▶ **LFU**: Pages waste buffer frames probably during the whole running time of the DB

## But ...

- ▶ Page reference pattern containing a loop slightly too long to fit in the buffer pool:
  - ▶ **OPT**: Hit rate close to 1
  - ▶ **LRU**: Hit rate of 0
- ▶ Some pages get referenced very frequently for a limited time:
  - ▶ **OPT**: Pages would be evicted after their last reference
  - ▶ **LFU**: Pages waste buffer frames probably during the whole running time of the DB
- ▶ Huge access time gap  $\implies$  Every saved page miss significantly improves the performance

## But ...

- ▶ Page reference pattern containing a loop slightly too long to fit in the buffer pool:
  - ▶ **OPT**: Hit rate close to 1
  - ▶ **LRU**: Hit rate of 0
- ▶ Some pages get referenced very frequently for a limited time:
  - ▶ **OPT**: Pages would be evicted after their last reference
  - ▶ **LFU**: Pages waste buffer frames probably during the whole running time of the DB
- ▶ Huge access time gap  $\implies$  Every saved page miss significantly improves the performance
- ▶ Pointer swizzling even amplifies that effect

## Subsection 1

### Probable pitfalls when Implementing a Page Eviction Strategy for a DBMS Buffer Manager

# General Problems Concerning DBMS Buffer Managers

# General Problems Concerning DBMS Buffer Managers

- Fixed pages cannot be evicted but a long timespan between a fix and an unfix of a page could make it a candidate for eviction.

# General Problems Concerning DBMS Buffer Managers

- ▶ Fixed pages cannot be evicted but a long timespan between a fix and an unfix of a page could make it a candidate for eviction.
- ▶ A page pinned for refix cannot be evicted but a long timespan in which a page is pinned could make it a candidate for eviction.



# General Problems Concerning DBMS Buffer Managers

- ▶ Fixed pages cannot be evicted but a long timespan between a fix and an unfix of a page could make it a candidate for eviction.
- ▶ A page pinned for refix cannot be evicted but a long timespan in which a page is pinned could make it a candidate for eviction.
- ▶ Dirty pages cannot be evicted but a page being dirty for a long timespan due to the update propagation using write-back policy could make it a candidate for eviction.

# Additional Problem When Using Pointer Swizzling

## Additional Problem When Using Pointer Swizzling

- ▶ A page containing swizzled pointer cannot be evicted but a page unfixed before the last unfix of one of its child pages could make it a candidate for eviction before its child pages got evicted.

# Solutions

# Solutions

- ▶ Check each of the restrictions before the eviction of a page.

# Solutions

- ▶ Check each of the restrictions before the eviction of a page.
- ▶ Update the statistics of the eviction strategy during an unfix, too.

# Solutions

- ▶ Check each of the restrictions before the eviction of a page.
- ▶ Update the statistics of the eviction strategy during an unfix, too.
- ▶ Update the statistics of the eviction strategy during an pin and unpin, too.

## Solutions

- ▶ Check each of the restrictions before the eviction of a page.
- ▶ Update the statistics of the eviction strategy during an unfix, too.
- ▶ Update the statistics of the eviction strategy during an pin and unpin, too.
- ▶ Use write-thru for update propagation or a page cleaner decoupled from the buffer pool as proposed in [SHG16].



# Solutions

- ▶ Check each of the restrictions before the eviction of a page.
- ▶ Update the statistics of the eviction strategy during an unfix, too.
- ▶ Update the statistics of the eviction strategy during an pin and unpin, too.
- ▶ Use write-thru for update propagation or a page cleaner decoupled from the buffer pool as proposed in [SHG16].
- ▶ Use a page eviction strategy that takes into account the content of pages (like the structure of an B tree).

## Subsection 2

### Evaluated Page Replacement Strategies

# RANDOM

## Overview

# RANDOM

## Overview

- ▶ Simplest page eviction strategy

# RANDOM

## Overview

- ▶ Simplest page eviction strategy
- ▶ Evicts a random page that can be evicted

# RANDOM

## Overview

- ▶ Simplest page eviction strategy
- ▶ Evicts a random page that can be evicted
- ▶ Won't evict frequently used pages as they're latched all the time

# GCLOCK

## Overview

# GCLOCK

## Overview

- ▶ Slight enhancement of the CLOCK algorithm: *generalized CLOCK*



# GCLOCK

## Overview

- ▶ Slight enhancement of the CLOCK algorithm: *generalized CLOCK*
- ▶ Uses finer-grained statistics about the recency of page references

# GCLOCK

## Overview

- ▶ Slight enhancement of the CLOCK algorithm: *generalized CLOCK*
- ▶ Uses finer-grained statistics about the recency of page references
- ▶ Parameter  $k$  defines granulation of statistics

# GCLOCK

## Overview

- ▶ Slight enhancement of the CLOCK algorithm: *generalized CLOCK*
- ▶ Uses finer-grained statistics about the recency of page references
- ▶ Parameter  $k$  defines granulation of statistics
  - ▶  $k = 1$ : CLOCK


# GCLOCK

## Overview

- ▶ Slight enhancement of the CLOCK algorithm: *generalized CLOCK*
- ▶ Uses finer-grained statistics about the recency of page references
- ▶ Parameter  $k$  defines granulation of statistics
  - ▶  $k = 1$ : CLOCK
  - ▶  $k = \#frames$ : Similar to LRU

# GCLOCK

## Example



# GCLOCK

## Advantage of Higher $k$ -Values

# GCLOCK

Advantage of Higher  $k$ -Values

Advantages of Lower  $k$ -Values

# GCLOCK

## Advantage of Higher $k$ -Values

- ▶ More detailed statistics about page references
  - ⇒ Higher hit rate
  - ⇒ Higher performance

## Advantages of Lower $k$ -Values



# GCLOCK

## Advantage of Higher $k$ -Values

- ▶ More detailed statistics about page references
  - ⇒ Higher hit rate
  - ⇒ Higher performance

## Advantages of Lower $k$ -Values

- ▶ Lower processing time required to find an eviction victim
  - ⇒ Higher performance

# GCLOCK

## Advantage of Higher $k$ -Values

- ▶ More detailed statistics about page references
  - ⇒ Higher hit rate
  - ⇒ Higher performance

## Advantages of Lower $k$ -Values

- ▶ Lower processing time required to find an eviction victim
  - ⇒ Higher performance
- ▶ Lower memory overhead due to shorter referenced-numbers

# GCLOCK

## Advantage of Higher $k$ -Values

- ▶ More detailed statistics about page references
  - ⇒ Higher hit rate
  - ⇒ Higher performance

## Advantages of Lower $k$ -Values

- ▶ Lower processing time required to find an eviction victim
  - ⇒ Higher performance
- ▶ Lower memory overhead due to shorter referenced-numbers

⇒ Trade-off between CPU- and I/O-optimization

# CAR

## Overview

# CAR

## Overview

- ▶ Extensive enhancement of the CLOCK algorithm: *Clock with Adaptive Replacement* [BM04]

# CAR

## Overview

- ▶ Extensive enhancement of the CLOCK algorithm: *Clock with Adaptive Replacement* [BM04]
- ▶ Approximation of the ARC page eviction strategy

# CAR

## Overview

- ▶ Extensive enhancement of the CLOCK algorithm: *Clock with Adaptive Replacement* [BM04]
- ▶ Approximation of the ARC page eviction strategy
- ▶ Uses two clocks and two LRU-lists

# CAR

## Overview

- ▶ Extensive enhancement of the CLOCK algorithm: *Clock with Adaptive Replacement* [BM04]
- ▶ Approximation of the ARC page eviction strategy
- ▶ Uses two clocks and two LRU-lists
- ▶ Advantages of CAR compared to CLOCK:



# CAR

## Overview

- ▶ Extensive enhancement of the CLOCK algorithm: *Clock with Adaptive Replacement* [BM04]
- ▶ Approximation of the ARC page eviction strategy
- ▶ Uses two clocks and two LRU-lists
- ▶ Advantages of CAR compared to CLOCK:
  - ▶ Weighted consideration of reference recency and frequency


# CAR

## Overview

- ▶ Extensive enhancement of the CLOCK algorithm: *Clock with Adaptive Replacement* [BM04]
- ▶ Approximation of the ARC page eviction strategy
- ▶ Uses two clocks and two LRU-lists
- ▶ Advantages of CAR compared to CLOCK:
  - ▶ Weighted consideration of reference recency and frequency
  - ▶ Scan-resistance

# CAR


## Example




## Subsection 3

### Performance Evaluation


# Buffer Pool Without Pointer Swizzling (TPC-C)




## Buffer Pool Without Pointer Swizzling (TPC-C)



# Buffer Pool Without Pointer Swizzling (TPC-C)




## Buffer Pool Without Pointer Swizzling (TPC-C)







# Buffer Pool Without Pointer Swizzling (TPC-C)




# Buffer Pool Without Pointer Swizzling (TPC-C)




# Buffer Pool Without Pointer Swizzling (TPC-C)




# Buffer Pool With Pointer Swizzling (TPC-C)




## Buffer Pool With Pointer Swizzling (TPC-C)




## Buffer Pool With Pointer Swizzling (TPC-C)



## Buffer Pool With Pointer Swizzling (TPC-C)




# Buffer Pool With Pointer Swizzling (TPC-C)







## Buffer Pool With Pointer Swizzling (TPC-C)




# Buffer Pool With Pointer Swizzling (TPC-C)




# Operation Performance




# Operation Performance




# Operation Performance



# Operation Performance



# Operation Performance



## Subsection 4

## Conclusion



# Conclusion

## Performance

# Conclusion

## Performance

- ▶ CAR has a significantly higher hit rate than RANDOM or GCLOCK

# Conclusion

## Performance

- ▶ CAR has a significantly higher hit rate than RANDOM or GCLOCK
- ▶ The hit rate of GCLOCK isn't significantly higher than the one of RANDOM

# Conclusion

## Performance

- ▶ CAR has a significantly higher hit rate than RANDOM or GCLOCK
- ▶ The hit rate of GCLOCK isn't significantly higher than the one of RANDOM
- ▶ Major differences in hit rate are only for buffer pool sizes of  $\leq \frac{1}{10}$  of the database size

# Conclusion

## Performance

- ▶ CAR has a significantly higher hit rate than RANDOM or GCLOCK
- ▶ The hit rate of GCLOCK isn't significantly higher than the one of RANDOM
- ▶ Major differences in hit rate are only for buffer pool sizes of  $\leq \frac{1}{10}$  of the database size
- ▶ The computational effort spent to do CAR eviction is 27–58 times higher

# Conclusion

## Performance

- ▶ CAR has a significantly higher hit rate than RANDOM or GCLOCK
- ▶ The hit rate of GCLOCK isn't significantly higher than the one of RANDOM
- ▶ Major differences in hit rate are only for buffer pool sizes of  $\leq \frac{1}{10}$  of the database size
- ▶ The computational effort spent to do CAR eviction is 27–58 times higher
- ▶ The overall performance of CAR isn't better than the one of RANDOM or GCLOCK

# References I



Sorav Bansal and Dharmendra S. Modha. “CAR: Clock with Adaptive Replacement”. Mar. 31, 2004.



Wolfgang Effelsberg and Theo Härder. “Principles of Database Buffer Management”. Dec. 1984.



Goetz Graefe et al. “In-Memory Performance for Big Data”. Sept. 2014.



Theo Härder and Erhard Rahm. *Datenbanksysteme - Konzepte und Techniken der Implementierung*. 2001. ISBN: 978-3-642-62659-3.



J. Eliot B. Moss. “Working with Persistent Objects: To Swizzle or Not to Swizzle”. Aug. 1992.

## References II



Lucas Sauer Caetano Lersch, Theo Härder, and Goetz Graefe. “Update propagation strategies for high-performance OLTP”. Aug. 14, 2016.



Seth J. White and David J. DeWitt. “QuickStore: A High Performance Mapped Object Store”. Oct. 1995.



# Your Turn to Ask ...